

CO211 - Operating Systems

4th October 2019

Outline of the Course

- overview and introduction structure, case studies
- processes and threads abstractions that an OS uses to execute code
- inter-process communication (IPC) allows multiple processes to communicate with each other
- memory management allocation, abstraction for virtual memory, paging
- device management types, drivers
- disk management scheduling, caching, RAID
- file systems basic abstractions for storage and implementation
- security authentication, access control

Note that this follows a similar structure to most OS courses, and therefore we can reference content from other sources. *Operating Systems: Three Easy Pieces* is recommended, as it bridges between this course and the PintOS lab.

Overview

The general overview is that there is a system bus that interconnects different hardware components (including CPU and memory), and allows for communication between them.

The operating system provides abstractions for programs to use, meaning that they do not have to deal with the complex hardware. For example, a process abstraction expects an interface to the hardware, which allows programs to be used on different hardware. This means that the OS will need how to control the hardware with drivers. The operating system has the following goals;

(1) managing resources

The operating system must be able to expose the resources efficiently to the application, and also share these resources fairly. Some examples are;

- CPU (multiple cores) should decide what runs on each hardware thread
- memory cache, RAM
- I/O devices displays (GPUs), network interfaces
- internal devices clocks, timers, interrupt controllers
- persistent storage

OS uses both time and space multiplexing for sharing. An example for the former is how the effect of parallelism can be achieved with a single CPU core by splitting up the time allocated per process, and an example for the latter is splitting up memory for each process.

On the other hand, with allocation, the OS must also support simultaneous resource access (such as to disks, RAM, network etc.). Continuing from this, it must also offer mutual exclusion, thus protecting risky operations (such as file writing). Generally, the OS aims to protect against corruption.

Finally, the operating system must also handle storing data, and enforce access control.

(2) clean interfaces

The OS should hide away the hardware, and applications use the hardware through an interface provided by the operating system. We can think of this as a virtual machine abstraction on top of the bare machine - similar to how the JVM works (but at a lower layer).

(3) concurrency and non-determinism

The operating system must be able to deal with concurrency, for example overlapping I/O and computation. This is because I/O devices tend to be slower, and while the device is working on the task, it shouldn't prevent the CPU from doing other work. An operating system may switch activities at arbitrary times, and this must be done safely - by offering synchronisation primitives. It should also protect processes by giving each program its own space, thus preventing interference.

Similarly, the OS is fundamentally non-deterministic, as it needs to handle interrupts (such as the network card receiving a packet, user interrupts, etc).

Tutorial Questions

- 1) List the most important resources that must be managed by an operating system in the following settings;
 - (a) supercomputer
 - computation time primarily used for intensive computations
 - memory
 - (b) networked workstations connected to a server
 - bandwidth must handle packet processing and network traffic
 - (c) smartphone
 - energy limited power, can power off unused hardware
 - mobile network (including other communication technology)
 - other sensors issues of privacy, when to expose GPS etc.

As this highlights, some uses will need specially designed operating systems. We also have general-process OS, as it takes a large amount of effort to implement a new operating system.

- 2) What is the **kernel** of an operating system?

The part of the OS is always in memory, and runs in the privileged part of the CPU (user mode cannot access all functionality). Implements commonly used functions of the OS and has complete access to all hardware.

Kernel Design

- **monolithic kernel**

Consider it as one large program that has all the functionality that you want an OS to perform.

The kernel is a single executable with its own address space. There exists a **system call** interface that allows user mode applications to access the hardware. Software invokes functionality from the kernel by issuing system calls - the CPU must switch from user mode to kernel mode to support this. The kernel then executes some instruction on behalf of the application. Device drivers are part of the monolithic kernel.

advantages

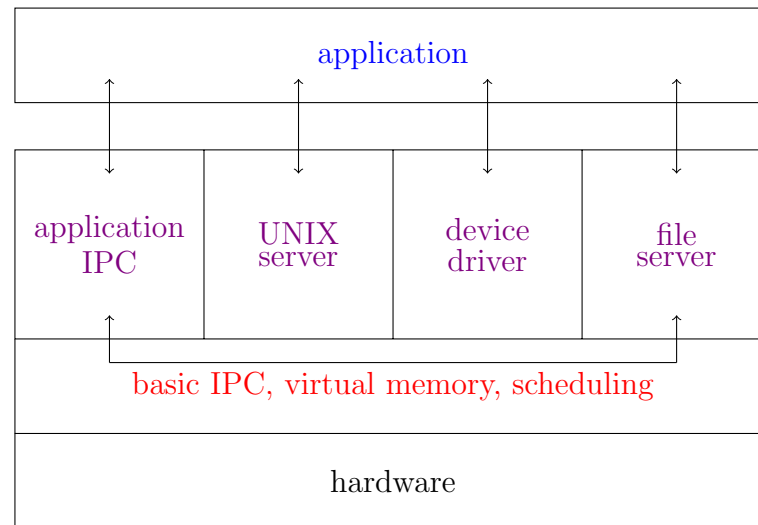
- efficient calls within the kernel, as there it remains in kernel mode
- flexible to write kernel components due to the shared memory (direct access with no limit to APIs)

disadvantages

- complex design
- no protection between bits of kernel functionality, therefore any bugs within the kernel will crash the entire machine

- **microkernels**

Only includes functionality that **requires** direct access to the hardware (or to be run in kernel mode). This is a minimalistic design and has the advantage of fewer bugs (due to the smaller amount of code).



Note that both the **application** and **servers** run in user mode, and the **kernel** is in kernel mode. The kernel performs IPC between the servers, which are separated for device I/O, scheduling, file access etc.

advantages

- less complex kernel
- clean interfaces for the servers
- more reliable; one of the servers could crash and then restart, without bringing the entire kernel down

disadvantages

- performance overhead due to the requirement of message passing and transitioning between user mode and kernel mode (checks must be done to maintain the separation) - less of an issue now due to better hardware (e.g. *Android*)

- **hybrid kernel**

many modern designs use a combination of both

This is a more structured design, however user-level servers can incur a performance penalty.

Linux Kernel

The structure of Linux system calls is to put arguments into registers Or on the stack, and then issue a trap to switch the CPU from user to kernel mode.

While C is the dominant language for the Linux kernel, the interrupt handlers are written in assembly, as they are low level pieces of code, and require fast performance (hence a low instruction count). Interrupt handlers are the primary means to interact with devices, it initiates dispatching which stops proxies, saves the state, starts the driver and returns.

Typically, we can split the Linux kernel into three parts;

- **I/O**

One of the design philosophies under UNIX style operating system is to treat everything as a file, and use this file abstraction to expose different resources. Therefore, a lot of I/O resources can be hidden under this virtual file system.

- **memory management**

Includes virtual memory with paging (and the abstractions associated with that).

- **process management**

Includes process and thread abstraction, as well as synchronisation and scheduling between them.

In addition to this, Linux supports dynamically loaded modules into the kernel. This support was important as it allowed for the hardware configuration to change (new device drivers could be loaded into the kernel, without recompiling).

Windows Kernel

The NTOS kernel layer implements Windows system call interface. This is an example of a hybrid kernel, as programs build on dynamic code libraries (DLLs) - which also make the kernel modular, however the executive servers in the kernel adopted the server model of the microkernel, but still runs in kernel space for the performance benefits. At the lower levels, there still exists a microkernel. In addition, there is also a hardware abstraction layer (HAL), as this was designed for portability.

It's also important to note that there are environment subsystems running in user mode allowing for different APIs to be exposed, including Win32, POSIX, and OS/2. While the Windows kernel was designed with a lot of flexibility, due to its nature as proprietary software, it only really focused support (until recently) on Win32 (and also Intel in terms of the HAL).