

Reinforcement Learning

(70028)

Markov Processes (Let's Go Markov)

In reinforcement learning, we need a real, tangible method for managing complexity. It's important to distinguish between **Markov Processes** and **Markov Decision Processes**.


A Markov process is a tuple $(\mathcal{S}, \mathcal{P})$, where \mathcal{S} is a set of states, and $\mathcal{P}_{ss'}$ is a matrix giving us the probability of transitioning from one state to another; note that the probabilities are based only on the current state (one at time t), and not looking beyond that - short memory is important;

$$\mathcal{P}_{ss'} = P[S_{t+1} = s' \mid S_t = s]$$

A Markov process generates a chain of states governed by probabilistic transitions.

A state s_t is Markov iff $P[s_{t+1} \mid s_t] = P[s_{t+1} \mid s_1, \dots, s_t]$; the conditional probability of transitioning to a particular state depends only on the particular state (previous states don't really matter). The equation states that the probabilities are equal, whether it be the current state, or all states preceding - the future is independent of the past given the present. Another way this can be thought of is that the present state, s_t , captures all information in the history of the agent's events, any data of the history is no longer needed once the state is known, or the current state is a sufficient statistic of the future.

An example is as follows, note that the black states are transient states (where it can lead to another state) and the **violet** states are terminal states. The following is Markovian as the probabilities don't change based on the history (how we got to a state).



The entries must be probabilities (hence between 0 and 1). The matrix defines transition probabilities from all states s to all successor states s' . Since all probabilities have to be accounted for (all rows of the matrix sum to 1) - after leaving s , we need to end up somewhere, which could also mean returning to s ;

$$\sum_{s'} \mathcal{P}_{ss'} = 1$$

In the example above, the probability of going to sleep after C2 (class 2) in the morning could be different depending on the time of day (i.e. constantly changing). If $P[s_{t+1} \mid s_t]$ doesn't depend on t , but rather just the origin and destination states, then the Markov chain is stationary or homogenous.

Markov Reward Process

A Markov Reward Process is a Markov chain which emits rewards (the reward hypothesis states that all of what we think of as goals and purposes can be thought of as the maximisation of the expected value of the cumulative sum of a scalar signal known as reward); hence a tuple $(\mathcal{S}, \mathcal{P}, \mathcal{R}, \gamma)$. This has the following components;


- \mathcal{S} a set of states
- $\mathcal{P}_{ss'}$ a state transition probability matrix
- $\mathcal{R}_s = \mathbb{E}[r_{t+1} | S_t = s]$
an expected immediate reward, collected upon departing state s (collection occurs at time $t + 1$, we are at state s at time t)
- $\gamma \in [0, 1]$ discount factor

We can define the return R_t as the total discounted reward from time-step t (note that we use $t + 1$ as the first element, since it's collected at $t + 1$);

$$R_t = r_{t+1} + \gamma r_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

The factor γ is how we discount the present value of future rewards; the value of receiving a reward r after $k + 1$ time steps is $\gamma^k r$, valuing immediate reward higher than a delayed reward - hence γ closer to 0 leads to short-sighted evaluation, whereas γ closer to 1 leads to far-sighted evaluation (taking future rewards more strongly).

We can add a reward to the previous example as follows (in red);



For example, consider a certain run, where the starting state $S_1 = C1$ and $\gamma = \frac{1}{2}$, T is the time to reach the terminal state;

$$R_1 = r_2 + \gamma r_3 + \dots + \gamma^{T-2} r_T$$

Consider the run where the student attends all classes in order and passes; hence C1, C2, C3, pass, sleep;

$$R_1 = -2 + \frac{1}{2} \cdot -2 + \frac{1}{2}^2 \cdot -2 + \frac{1}{2}^3 \cdot 10$$

Most MRPs are discounted with $\gamma < 1$, as it's mathematically convenient by avoiding infinite returns in cyclic / infinite processes (by causing convergence). It also aids in expressing uncertainty in future

rewards. A more tangible example is a financial reward, where immediate rewards can be put into a bank and earn interest, similarly, animal decision making shows preference for immediate rewards rather than future rewards.

We can define the state value function $v(s)$ of a MRP as the expected return R starting from state s at time t , thinking of the state as a function parameter;

$$v(s) = \mathbb{E}[R_t \mid S_t = s]$$

The lecture then goes over an example using golf, which is actually quite intuitive.

The Bellman Equation for MRPs is as follows. We can express it in a recurrence relation, as the **immediate reward** and the **discounted return of the successor state**.

$$\begin{aligned} v(s) &= \mathbb{E}[R_t \mid S_t = s] \\ &= \mathbb{E}[r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \mid S_t = s] \\ &= \mathbb{E}[r_{t+1} + \gamma(r_{t+2} + \gamma r_{t+3} + \dots) \mid S_t = s] \\ &= \mathbb{E}[r_{t+1} + \gamma R_{t+1} \mid S_t = s] \\ &= \mathbb{E}[r_{t+1} + \gamma v(S_{t+1}) \mid S_t = s] \end{aligned}$$

The equation can also be written as the sum notation (the previous one was the expectation notation, this has the expectation written out) - there are a total of n of these equations, as there's one for each state;

$$v(s) = \mathcal{R}_s + \gamma \sum_{s' \in S} \mathcal{P}_{ss'} v(s')$$

As such, this can be written in vector notation as follows, with \mathbf{v} being n -dimensional;

$$\mathbf{v} = \mathcal{R} + \gamma \mathcal{P} \mathbf{v}$$

This can be directly solved as follows, as it's linear and self-consistent;

$$\begin{aligned} \mathbf{v} &= \mathcal{R} + \gamma \mathcal{P} \mathbf{v} \\ (\mathbf{1} - \gamma \mathcal{P}) \mathbf{v} &= \mathcal{R} \\ \mathbf{v} &= (\mathbf{1} - \gamma \mathcal{P})^{-1} \mathcal{R} \end{aligned}$$

Since matrix inversion is computationally expensive, being in the order of n^3 for n states, a direct solution is only feasible for small MRPs. Iterative methods for solving large MRPs include (and all three will be covered);

- dynamic programming
- Monte-Carlo evaluation
- Temporal-Difference learning

Policies

A policy π is a function of the state, formalising the actions to take at a given state. A rigid, deterministic policy can be disadvantageous (e.g. rock, paper, scissors) - exposing the agent to being systematically exploited. A policy can be formally described as the conditional probability distribution to execute an action $a \in \mathcal{A}$ given that one is in state $s \in \mathcal{S}$ at time t ;

$$\pi_t(a, s) = P[A_t = a \mid S_t = s]$$

The general form of the policy is probability, or stochastic, hence π is a probability. However, if the policy is deterministic (only a single a is possible for state s), then $\pi(a, s) = 1$, $\pi(a', s) = 0$, $\forall a \neq a'$.

Consider the following example, where there are two actions, a_1, a_2 where we either play the lottery (costing 1), or save (not costing anything). The two states, s_1 and s_2 correspond to winning or losing the lottery.


$$a^* = \operatorname{argmax}_{a_i} \sum_{j=1}^2 \mathcal{R}_{s_j}^{a_j} P[s_j \mid a_i]$$

Markov Decision Process

The emphasis decision process, with decision being the key, combines the policies with MRPs. The MDP consists of the following;

- \mathcal{S} state space
- \mathcal{A} action space
- $\mathcal{P}_{ss'}^a$ transition probability $p(s_{t+1} | s_t, a_t)$
probability of transitioning to the next state s_{t+1} , given the current state s_t and action a_t taken
- $\gamma \in [0, 1]$ discount factor
- $\mathcal{R}_{ss'}^a = r(s, a, s')$ immediate reward function
in temporal notation, $r_{t+1} = r(s_{t+1}, s_t, a_t)$ - reward is collected upon the transition from s_t to s_{t+1} , which occurs at time $t + 1$
- π policy, can be either stochastic or deterministic
stochastic is written as the following; $\mathbf{a} \sim p_\pi(\mathbf{a} | \mathbf{s}) = \pi(\mathbf{a} | \mathbf{s}) \equiv \pi(a, s)$ - being a probability distribution
deterministic is written as $\mathbf{a} = \pi(\mathbf{s})$ (indicator function)

Note that the transition probability and policy both take the action into account, as parameters. We can graphically represent this as follows, with nodes denoting variables, and edges denoting conditional dependencies between these variables;



This tells us that the action a_t depends on the current state s_t , the next state s_{t+1} depends on both the action and the current state, and the reward r_{t+1} depends on all three.

Value Function

The goodness of a given state is defined with the value function (where R_t is a discounted total return, and r_{t+k+1} are immediate rewards);


$$V^\pi(s) = \mathbb{E}_\pi[R_t | S_t = s] = E \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid S_t = s \right]$$

This is quite similar to the derivation of the Bellman equation for MRPs, but now including the action (see the policies π). Note that expectation is a linear operator, hence we can justify the final line, also note in the penultimate line we separate out the **next reward** from the discounted rewards;

$$\begin{aligned}
 V^\pi(s) &= \mathbb{E}_\pi[R_t | S_t = s] \\
 &= \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid S_t = s \right] \\
 &= \mathbb{E}_\pi \left[\textcolor{violet}{r}_{t+1} + \gamma \sum_{k=0}^{\infty} \gamma^k r_{t+k+2} \mid S_t = s \right] \\
 &= \mathbb{E}[r_{t+1} | S_t = s] + \gamma \mathbb{E} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+2} \mid S_t = s \right]
 \end{aligned}$$

Backup Diagrams

We start at a white node, at a particular state s (since we are conditioning on a particular state $S_t = s$). From this state, we can take several actions, represented by the black nodes, which leads us to following states s' , with a reward r . This state value information is transferred back up to s from its successor state s' , performing the **update** or **backup** operation at the heart of the reinforcement learning method.



In order to calculate the value of state s , we need to average over all possible traces, which is what's going on behind the scenes in the expectation operator - an average weighted by probabilities. All of which live inside the MDP;

- probability of the chosen action a is given by the policy $P[a | s] = \pi(a, s)$
- probability of a transition to s' is given by the transition probability $P[s' | s, a] = \mathcal{P}_{ss'}^a$
- instantaneous reward r is given by the reward function $r(s, a, s') = \mathcal{R}_{ss'}^a$
- the value of the next state s' , weighted by the probability functions is given recursively by $v(s')$

Writing it out, note that we have the value function of the state s' in **violet**;

$$\begin{aligned}
 \mathbb{E}[r_{t+1} | S_t = s] &= \sum_{a \in \mathcal{A}} P[a | s] \left(\sum_{s' \in \mathcal{S}} P[s' | s, a] r(s, a, s') \right) \\
 \mathbb{E} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+2} \middle| S_t = s \right] &= \sum_{a \in \mathcal{A}} P[a | s] \left(\sum_{s' \in \mathcal{S}} P[s' | s, a] \sum_{k=0}^{\infty} \gamma^k r_{t+k+2} \right) \\
 V^\pi(s') &= \mathbb{E}[R_{t+1} | S_{t+1} = s'] \\
 &= \mathbb{E} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+2} \middle| S_{t+1} = s' \right]
 \end{aligned}$$

We can combine all of this as follows;

$$V^\pi(s) = \mathbb{E}_\pi[R_t | S_t = s] \quad (1)$$

$$= \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \middle| S_t = s \right] \quad (2)$$

$$= \mathbb{E}_\pi \left[r_{t+1} + \gamma \sum_{k=0}^{\infty} \gamma^k r_{t+k+2} \middle| S_t = s \right] \quad (3)$$

$$= \sum_{a \in \mathcal{A}} \pi(a, s) \left(\sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \left(\mathcal{R}_{ss'}^a + \gamma \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+2} \middle| S_{t+1} = s' \right] \right) \right) \quad (4)$$

$$= \sum_{a \in \mathcal{A}} \pi(a, s) \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a (\mathcal{R}_{ss'}^a + \gamma V^\pi(s')) \quad (5)$$

Here we are performing the following steps;

- (2) write the definition of the return
- (3) separate immediate reward

- (4) split expectation in two, as it's a linear operator, also write out expectation weighted by probabilities, and using proper notation for policies $\pi(a, s)$, transition probabilities $\mathcal{P}_{ss'}^a$, and reward function $\mathcal{R}_{ss'}^a$
- (5) substitute with recursive definition


This is a consistency condition imposed on the value function and also has a unique solution. Computing the value function for an arbitrary policy is known as policy evaluation or prediction problem. Now, we need to iterate applications to obtain better estimates - note that the subscripts in $V_1(s), V_2(s), \dots, V_k(s)$ denote iterations, not states; this is guaranteed to converge. This is known as iterative policy evaluation. A stopping condition can be achieved by checking that the **largest** change in the value function, between iterations, is below a certain small threshold. This can be formalised as follows - note that the value function is on a particular policy;

1. input π , the policy to be evaluated
2. initialise $V(s) = 0$ for all $s \in \mathcal{S}^+$
3. repeat the following until $\Delta < \theta$ (where θ is some small positive number)
 - (a) $\Delta \leftarrow 0$
 - (b) for each $s \in \mathcal{S}$;
 - i. $v \leftarrow V(s)$ store old value
 - ii. $V(s) \leftarrow \sum_{a \in \mathcal{A}} \pi(a, s) \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a (\mathcal{R}_{ss'}^a + \gamma V^\pi(s'))$ sweep through successors, a full backup
 - iii. $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
4. output $V \approx V^\pi$

Note that this replaces values, in place, converging faster than a two-array method, which would have both an old and new array.

Stair Climbing MDP

Consider the following example; for brevity, a **violet** edge means a Right action, and a **teal** edge denotes a Left action. P and G are both absorbing / terminal states, assume we start at s_3 with $\gamma = 0.9$, an unbiased policy (such that all actions are equally probable) with $\pi(s, L) = \pi(s, R) = 0.5$, hence randomly selecting actions.



Note that in the first iteration, the only changes are to s_1 and s_5 , as they are the only ones with successor states that have different rewards (all other states will cancel out), also note the symmetry stemming from the symmetrical problem.

V	P	s_1	s_2	s_3	s_4	s_5	G
V_0	0	0	0	0	0	0	0
V_1	0	-5.5	0	0	0	5.5	0
V_2	0	-5.5	-2.48	0	2.48	5.5	0
V_2	0	-6.61	-2.48	0	2.48	6.61	0
\vdots							
V_∞	0	-6.9	-3.1	0	3.1	6.9	0

Note that we are still equally likely to go to the left, despite being significantly worse; thus knowing the value of the policy can improve the policy.

State-Action Value Function

This takes in two parameters; a state s and an action a , giving us a function that determines the value of taking a certain action at a state.

$$Q^\pi(s, a) = \mathbb{E}[R_t \mid S_t = s, A_t = a] = \mathbb{E} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid S_t = s, A_t = a \right]$$

The relation between the state value function and this is;

$$V^\pi(s) = \sum_{a \in \mathcal{A}} \pi(s, a) Q^\pi(s, a)$$

Bellman Optimality Equations

Previously, we discussed arbitrary policies. However, we can define an ordering on policies (such that some policies are better than others) by saying a policy is better than, or equal to, another policy if its expected return is also greater than or equal to the other policy for all states; $\pi \geq \pi'$ iff $\forall s \in \mathcal{S} [V^\pi(s) \geq V^{\pi'}(s)]$. As such, the optimal value function is defined as;


$$V^*(s) = \max_{\pi} V^\pi(s), \forall s \in \mathcal{S}$$

We call the policy π^* which maximises the value function the optimal policy; there will always be at least one, but multiple can exist. Similarly, there is also an optimal state-action value function;


$$Q^*(s, a) = \max_{\pi} Q^\pi(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A} = \mathbb{E}[r_{t+1} + \gamma V^*(s_{t+1}) \mid S_t = s, A_t = a]$$

The Bellman equations for these are called Bellman Optimality equations.

We have already seen the backup diagram for the value function, and the state-action value function backup diagram is similar (you can think of them as the black nodes and its children);



The white nodes are associated with the value function $V^\pi(s)$, the black nodes with the value-action function $Q^\pi(s, a)$, and the paths between the nodes taken with probability $\pi(s, a)$. The relationship for the function, on just the value function backup diagram can be shown as follows;



If we want the optimal value for a state, only actions that give the highest value should be chosen;

$$V^*(s) = \max_{a \in \mathcal{A}} \sum_{a \in \mathcal{A}} \pi(s, a) Q^\pi(s, a) \quad (1)$$

$$= \max_a Q^{\pi^*}(s, a) \quad (2)$$

$$= \max_a \mathbb{E}[R_t \mid S_t = s, A_t = a] \quad (3)$$

$$= \max_a \mathbb{E} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid S_t = s, A_t = a \right] \quad (4)$$

$$= \max_a \mathbb{E} \left[r_{t+1} + \gamma \sum_{k=0}^{\infty} \gamma^k r_{t+k+2} \mid S_t = s, A_t = a \right] \quad (5)$$

$$= \max_a \mathbb{E}[r_{t+1} + \gamma V^*(s_{t+1}) \mid S_t = s, A_t = a] \quad (6)$$

$$= \max_a \sum_{s'} P[s' \mid s, a] (r(s, a, s') + \gamma V^*(s')) \quad (7)$$

$$= \max_a \sum_{s'} \mathcal{P}_{ss'}^a (\mathcal{R}_{ss'}^a + \gamma V^*(s')) \quad (8)$$

We perform the following steps;

- (3) write down definition of state-action value function, being the expected return conditioned on s and a
- (4) perform usual expansion with the maximum on the left
- (7) replace with probabilities

There is no reference to any particular policy and must therefore be satisfied by all optimal policies. The optimality equation expresses that the value of a state under an optimal policy is equal to the expected return of the best action from the state. This can be done similarly for Q^* , except the maximum is now on the inside;

$$\begin{aligned} Q^*(s, a) &= \mathbb{E} \left[r_{t+1} + \gamma \max_{a'} Q^*(s_{t+1}, a') \mid S_t = s, A_t = a \right] \\ &= \sum_{s'} \mathcal{P}_{ss'}^a \left(\mathcal{R}_{ss'}^a + \gamma \max_{a'} Q^*(s', a') \right) \end{aligned}$$

Notice that the equation doesn't require π^* at all, which is useful as we don't need to know the optimal policy to solve the optimality equations. For finite MDPs, this equation has a unique solution, independent of the policy. The bellman optimality equation is a set of N non-linear equations, where $N = |\mathcal{S}|$, with N unknowns.

An explicit solution for the optimality equation provides one route for an optimal policy - however we are often going to encounter a high-dimensional problem (large state space). This also assumes the following, which are rarely true;

- we accurately know the dynamics of the environment
- we have the resources to find the solution
- the Markov property

We therefore often settle for approximate solutions.

The BOE convergence theorem states that for an MDP with finite state and action space, the optimality equations have a unique solution and the values produced by iteration converge to the solution of the equations. The proof of this rests on the Banach Fixed Point / Contraction Mapping Theorem.